



Learning Quadruped Locomotion from Casual Videos — additional experiments: all figures

A — Exp 1: global keypoint trajectories, MoCap vs video (speed-matched right turns)

— MoCap (right turn, 1.47 s, 60 Hz) — Video w. Depth Camera (right turn, 1.47 s, 30 Hz)




J — Base-centre paths, all clips origin-aligned to a common heading




Base centre only, top-down. Every clip is translated to the origin (square) and rotated so its initial heading points along +x, so only the shape of the motion differs. Dots mark each clip's end; red = clips added by the extended dataset.


C — Exp 2: command-space coverage, MoCap / Video / Video (extended)




F — Exp 2: coverage of the 24-dim AMP discriminator state




E — Exp 3: per-clip hand-coding in the pipeline




I — Global motion paths of every clip (world coordinates, metres)




B — Exp 1: noise metrics + comparison across the four depth sources



G — Exp 2: target command distribution vs demonstrated velocity density



D — Exp 2: terrain / posture coverage (stairs, box, stand-up)



H — Diagnostics: described filter vs hand-coding, and the intrinsic bug

