



— MoCap (right turn, 1.47 s, 60 Hz) — Video w. Depth Camera (right turn, 1.47 s, 30 Hz)


(a) global 3D keypoint trajectories




(b) rigid-torso violation





(c) top-down projection





(d) sagittal projection





MoCap
(6 clips)



Video flat
(8 clips)




Video stairs/box/
stand-up (6 clips)




One rigid transform per clip (frame-0 rear keypoint to origin, frame-0 heading to +x); nothing re-centred per frame.
Thick = base path, thin = the four feet. Dots mark each clip's start. Colour distinguishes clips within a row.


**MoCap
(6 clips)**




**Video
(4 clips)**



**Video (extended)
(8 clips)**




**Video: stairs / box / stand-up
(6 clips)**




Base centre only, top-down. Every clip is translated to the origin (square) and rotated so its initial heading points along +x, so only the shape of the motion differs. Dots mark each clip's end; red = clips added by the extended dataset.


M1 torso-length CV
[%]




M2 SG residual
[% of L]




M4 normalized jerk
[1]






M2 jitter by source
[% of L]






Overlapping feet
[% of frames]






MoCap



Video




Video (extended)




not covered MoCap only new in video covered by both


MoCap




Video




Video (extended)




MoCap







Video



Video (extended)




not covered MoCap only new in video covered by both

(a) posture**(b) dynamics****(c) clearance ECDF****(d) per clip**


MoCap
Video flat

Video box
Video stairs/stand-up


(a) per-clip hand-coding in the pipeline




(b) normalized cost



(a) slow_turn_1771233000, keypoint VL



(b) recall per clip



(c) camera-calibration bug

